

English

Español

 Download PDF

houston

Bootcamp · From 0 to 1 with AI agents

• STEP-BY-STEP GUIDE · LIVE BOOTCAMP

From 0 to 1 with AI agents

Everything we did in the bootcamp, in order and at your own pace, with the slides from the session. You finish with an agent installed, connected to your tools and working on a real case: pulling a prospect database out of a LinkedIn post.


 2 hours at your own pace For business owners and operators Mac, Windows or Linux No coding

houston is not a bootcamp exercise: it's a real product and you can put it to work in your company. At the end of the guide I show you [how we do this with teams](#) and how to join the beta program. First, build yours.

WHAT YOU'LL WALK AWAY WITH

An agent that goes out to the internet and builds your customer list

This isn't a demo. It's an agent that stays installed in your account, with its instructions, its skill and its memory, and keeps working after you close this guide.



WHERE IT ALL STARTS The houston **Agent Store**. That's where you get the Sales Rep agent, already trained, in two clicks.

You'll understand it

The method

What an agent is made of, and how to go from a vague business problem to a specific agent, using the canvas we filled in live.

This is the part that serves you forever. Today's agent solves one case; the method solves all of yours.

You get it prebuilt

The Sales Rep agent

You hand it the link to a LinkedIn post with comments and it hands you back a Google Sheet with the name, role, company and comment of every person who commented.

It comes with instructions, a skill and memories already written. Its full instructions are at the end of the guide.

The tools you'll connect



Apify · eyes on the web



OpenRouter · 400+ models



Google Sheets · the database



LinkedIn · the source



Anthropic



OpenAI

The result, concretely

This is how the sheet the agent builds looks. Every person who commented, with their role, their company, what they wrote and their profile.



[Prospects](#) · [Agents bootcamp](#)

[Multiplayer](#)
[Agents](#)
[Features](#)
[Pricing](#)
[Questions](#)

docs.google.com/spreadsheets

Name	Role	Company	Comment	Pro...	Email
Ana M.	Head of Growth	Nubetech	"Houston"	in/a...	ana@nubetech.co
Carlos R.	Founder & CEO	Veloro	"Houston 🚀"	in/c...	carlos@veloro.io
Diana P.	VP Sales	Aurora Labs	"Houston"	in/di...	diana@auroralab...
Javier T.	Product Lead	Kantto	"Houston!"	in/ja...	not found
Mariana B.	Marketing Director	Origen	"Houston 🗯️"	in/m...	mariana@origen....
Sergio D.	COO	Brío	"Houston"	in/s...	searching...

Prospects

Enriched

Sequence

EXAMPLE The sheet fills itself while the agent works. An empty email cell is normal: nobody finds 100%.

💡 **THE POWER MOVE**

It doesn't have to be your post. Run the same recipe on a post from a leader in your industry, from a competitor or from an event, and **their audience becomes your list**. A post with a thousand comments turns into a database of a thousand prospects in minutes, not days.

⚠️ **USE IT RESPONSIBLY**

This is real power. The agent **never sends anything on its own**: it leaves everything waiting for your approval. Keep that line in place and prospect with judgment.

📄 **ONE IDEA FOR THE WHOLE RIDE**

An agent is **a new intern you're training**. It's already smart: what it doesn't know is how *you* work. Everything you already know about training someone new applies here.



It's the WhatsApp group of everyone who was in the session. We unblock whatever gets stuck, share the codes and the new tutorials, and everyone tells what agent they built. The group speaks mostly Spanish, and English is welcome. Scan the code with your phone or use the button.

Join the community ↗

chat.whatsapp.com/BdLpn8T92DL4pRahiqOZEer

The 5 pieces of an agent

By the end of this guide you'll have touched all five. Come back to this slide whenever something doesn't click.

ANATOMÍA

Todo agente está compuesto de cinco cosas:

01

Identidad

QUIÉN ES

Rol, voz y barreras operativas. Define de qué es responsable el agente y cómo te representa.

"Tú eres nuestro SDR. Tono: marca. No puedes enviar correos externos sin aprobación."

02

Habilidades

QUÉ SABE HACER

Procedimientos componibles que el agente puede ejecutar. Se construyen una vez y se reutilizan entre agentes y equipos.

"Secuencia de prospección. Calificación de leads. Escalamiento por Slack."

03

Herramientas

QUÉ USA

Sistemas externos de los que el agente lee y en los que escribe. Las apps que tu equipo ya usa todos los días.

04

Memoria

QUÉ RECUERDA

Contexto persistente entre sesiones. Decisiones previas, terminología, preferencias y relaciones.

"Conoce a tus mejores clientes. Sabe qué significa 'ola Q3:'"

EL CEREBRO
El modelo de IA

La quinta pieza: el modelo que razona y decide. Con Houston eliges el mejor para cada tarea.

OpenAI

Anthropic

Gemini

+400

Houston · Agentes de IA que realmente hacen el trabajo

gethouston.ai

09

SLIDE 09 Identity, skills, tools and memory. The fifth piece, the brain, is the AI model you choose. The slides are the originals from the live session, in Spanish; the captions carry the point in English.

⚠ THE GOLDEN RULE

A new agent **doesn't send anything on its own**. First it proposes, you review and approve. Trust is earned with reviewed results, exactly like with a person.



All 49 slides from the session as a PDF, exactly as they were shown live. Handy for reviewing without the guide open, or for walking your team through it. The deck is in Spanish and weighs 11 MB.

Download the PDF 

0

STEP 0

Before you start

~10 minutes

Four things within reach. All of them are free and take minutes.

A computer

Mac, Windows or Linux, charged and with Chrome open. houston doesn't run on a phone or a tablet.

A Google account

So the agent can create and fill the Google Sheet with your prospects.

An AI account

ChatGPT (OpenAI) or Claude (Anthropic). Either one gets you started. OpenRouter works too, and you'll set it up in Step 4.

A LinkedIn post

One with lots of comments. It can be yours or from anyone with a following in your industry.

The three bootcamp codes

You'll use them in steps 3 and 4. Copy them now.



houston invite · on the download page at gethouston.ai

getsh*tdone



Apify coupon · 75 dollars in credit, valid for two months

HOUSTON_ROCKET



OpenRouter code · 10 dollars in credit, all caps

OR-GETHOUSTON



Tres cosas para instalar. Vamos una por una, juntos.

01

Descarga Houston

La plataforma donde vive tu agente. Ahí lo creas y lo pones a trabajar.

ENTRA A

gethouston.ai

INVITACIÓN **getsh*tdone**

02

Crea tu cuenta en Apify

Le da a tu agente los ojos en la web: extrae datos de sitios sin API.

\$75 USD gratis para ti

ESCANEA

Crea tu cuenta

CUPÓN **HOUSTON_ROCKET**

03

Crea tu cuenta en OpenRouter

Le da el cerebro: una cuenta para usar cualquier modelo dentro de Houston.

\$10 USD gratis para ti

ESCANEA

Crea tu cuenta

CÓDIGO **OR-GETHOUSTON**

Nadie se queda atrás

Los links y los códigos están en el chat. Si te atascas, escríbelo y un helper te acompaña sin que paremos a todos.

Houston · Agentes de IA que realmente hacen el trabajo

gethouston.ai

34

SLIDE 34 The three workshop accounts, exactly as we saw them live. The codes are the same ones above.

Open them from your phone or your computer

Scan the code with your phone camera, or click the button if you're already at the computer. The coupons come pre-applied in the link.



Eyes on the web. 75 dollars in credit, valid for two months.

HOUSTON_ROCKET

Create account ↗



The brain: 400+ models with a single account. 10 dollars in credit.

OR-GETHOUSTON

Redeem credits ↗



houston

The app where your agent lives. Download it on your computer, not your phone.

getsh*tdone

Go to gethouston.ai ↗



Bootcamp community

The WhatsApp group from the session. If something gets stuck, that's where to ask.

Join the group ↗

**ARE YOU SET?**

You've got the computer handy, you can open your email and you have a quiet stretch ahead. You don't have to do it all in one sitting: every step is saved.

THE FULL ROUTE

From step 0 to 6, in order

Steps 1 and 2 are the method. From 3 to 6 it's all keyboard. Do them top to bottom.

0**Before you start**

Accounts, codes and a LinkedIn post

~10 min

1**What an agent is made of**

The 5 pieces, with the slides from the session

~20 min

2**Pick the right problem**

The MUUDS canvas, same as live

~20 min

3**Install houston**

Download, invite code and first launch

~15 min

4**Open Apify and OpenRouter**

Your agent's eyes and brain

~15 min

5**Install your agent**

Agent Store, search for "Felipe"

~10 min

1

STEP 1 OF 6

What an agent is made of

~20 minutes

Twenty minutes of reading that save you weeks of flailing around. Here are the slides from the session, in the same order. The bootcamp ran in Spanish, so the slides are in Spanish: every caption carries the point in English, and you lose nothing by not reading them. If you'd rather start with your hands, jump to [Step 3](#) and come back later.

THE COMPLETE SLIDES



Below are the ones that matter, but you can take them all

The entire bootcamp deck as a PDF: 49 slides, including the ones that didn't make it into this guide. In Spanish.

[Download the PDF](#) 

Chatbot, co-worker and agent are not the same thing

The industry calls everything an "agent". Here's what each one actually does:



La industria usa la palabra "Agentes" para todo. Esto es lo que realmente hace cada uno.



Chatbot

RESPONDE PREGUNTAS

Reactivo y sin memoria. Útil cuando ya sabes qué preguntar.

"¿Cuál es nuestra política de devoluciones?"

Houston · Agentes de IA que realmente hacen el trabajo



Co-worker

ASISTE AL HUMANO

Sugiere los siguientes pasos mientras el humano sigue al volante. Tú decides qué sale a producción.

"Aquí está un borrador de correo basado en los datos."

gethouston.ai



Agente

HACE EL TRABAJO

Razona, decide y ejecuta acciones a través de tus sistemas. Tú gestionas resultados, no pasos.

"Saca los datos, redacta el correo, lo envía a la persona correcta y registra la acción."

08

SLIDE 08 The chatbot answers, the co-worker assists, the agent does the work. Only the third one executes across your systems.

1 Chatbot · answers questions. Reactive and without memory. Useful when you already know what to ask.

2 Co-worker · assists the human. It suggests next steps, but you're still at the wheel.


3 Agent · does the work. It reasons, decides and executes. You manage results, not steps.

→ No tools, no agent. If it can't read and write in your apps, it's a chatbot with a good vocabulary.


The car analogy

If you take one single image from the whole bootcamp, make it this one.


El modelo de IA es el motor. Houston es el resto del carro.



- 1 **El volante** · Identidad: quién conduce y a dónde va.
- 2 **La palanca de cambios** · Habilidades: las maniobras que sabe ejecutar.
- 3 **Las ruedas** · Herramientas: donde el agente toca la calle, tus apps.
- 4 **El tablero y el GPS** · Memoria: el recorrido no se olvida.
- 5 **El motor** · El cerebro: el modelo que genera la potencia.



Houston · Agentes de IA que realmente hacen el trabajo



gethouston.ai

10

SLIDE 10 The AI model is the engine. houston is the rest of the car: steering wheel, gearshift, wheels and dashboard.



WHY IT MATTERS

When you use Claude or ChatGPT you're tied to that lab's engine. A car like houston doesn't belong to any of them: you swap the engine whenever you want, even from the chat.

Piece 1 · Identity: who it is

It lives in **Settings** → **Instructions** . Five ingredients: **role**, **behavior**, **hard rules**, **vocabulary** and **tone**. This is what makes two agents running the same model feel completely different.

El mismo agente. Dos instrucciones diferentes. Resultados completamente distintos.

MALAS INSTRUCCIONES

Agente de ventas

Eres un agente de ventas.

Ayúdame a conseguir más clientes.
Escribe correos de prospección.
Sé profesional.

Por qué falla: sin alcance, sin reglas y sin vocabulario propio. El agente inventa precios, manda correos que no aprobaste y cada día responde distinto.

Houston · Agentes de IA que realmente hacen el trabajo

BUENAS INSTRUCCIONES

Operador comercial de Houston

Rol
Eres mi operador comercial completo: prospección, pipeline, reuniones, propuestas y CRM.

Comportamiento
Redactas. Nunca envías.
Si te falta un dato, preguntas una sola vez.

Reglas duras
Nunca envías un correo sin mi aprobación.
Nunca mueves etapas del CRM ni cambias precios.
Si la fuente es débil, marcas TBD. No inventas.

Vocabulario
VERDE = asignar. AMARILLO = nutrir. ROJO = soltar.
Cliente ideal = empresas de 20 a 200 empleados.

Tono
Directo, sin tecnicismos. Sin nombres de archivos.

gethouston.ai

12

SLIDE 12 Same agent, two different sets of instructions, completely different results.

Bad instructions look like this. No scope, no rules and no vocabulary of its own: the agent invents prices, sends emails you never approved, and answers differently every day.

NOT LIKE THIS

You're a sales agent. Help me get more customers. Write prospecting emails. Be professional.

Good instructions look like this. Copy them and adapt them to your business: it's the template that works for any role.

COPY AND PASTE INTO YOUR AGENT'S INSTRUCTIONS

Commercial operator ## Role You are my full commercial operator: prospecting, pipeline, meetings, proposals and CRM. ## Behavior You draft. You never send. If you're missing a fact, you ask once. ## Hard rules You never send an email without my approval. You never move CRM stages or change prices. If the source is weak, you mark it TBD. You don't make things up. ## Vocabulary GREEN = assign. YELLOW = nurture. RED = drop. Ideal customer = companies with 20 to 200 employees. ## Tone Direct, no jargon. No file names.



WHAT ALMOST NOBODY WRITES DOWN

The **hard rules** block. Half of a good onboarding is what the agent must *not* do. It's the same with people, and we almost never put it in writing.

Piece 2 • Skills: what it knows how to do

A skill isn't anything esoteric: it's **a text file with instructions** that the agent reads only when it needs it. It lives in **Settings** → **Skills**. Written once, used thousands of times.

HABILIDADES - CÓMO SE ESCRIBE UNA

Una habilidad es un archivo con instrucciones. El agente lo lee solo cuando lo necesita.

LOS 3 NIVELES DE CARGA

01 **Nombre + descripción**
Siempre en contexto. Es lo único que el agente ve antes de decidir si usa la habilidad.

02 **El cuerpo del SKILL.md**
Se carga solo cuando la habilidad se activa. Tus reglas, tu proceso, tus ejemplos. Menos de 500 líneas.

03 **Recursos opcionales**
Scripts, plantillas y documentos de referencia. Se leen solo bajo demanda, sin gastar contexto.

Piénsala como la guía de inducción de un empleado brillante: ya es inteligente, solo no conoce tu forma de trabajar.

ASÍ SE VE POR DENTRO

```
live-linkedin-comments-to-sheets/SKILL.md

---
name: live-linkedin-comments-to-sheets
description: Dada la URL de un post de LinkedIn,
crea una hoja, envía el link al chat de inmediato
y la va llenando con cada comentarista, en vivo.
Úsala cuando Felipe diga "prospecta a esta
publicación de LinkedIn".

---

# Comentarios de LinkedIn a una hoja, en vivo

## La regla dura
En cuanto la hoja exista, envía el link al chat y
sigue trabajando. Nunca esperes al scraper.

## Reglas
- Columnas: Nombre | Cargo | Empresa | Comentario | Email
- Escribe por lotes: el público ve las filas llegar.
- Ordena A-Z una sola vez, al final de todo.

## Cuando NO usarla
- Extracciones sin público: usa otra habilidad.
```

Houston • Agentes de IA que realmente hacen el trabajo

gethouston.ai

14

SLIDE 14 On the right, the real skill the Sales Rep agent uses: name, description, the hard rule and when not to use it.

1 Name and description. Always in context. It's the only thing the agent sees before deciding whether to use the skill.

2 The body. Loaded only when triggered. Your rules, your process, your examples. Under 500 lines.

3 Optional resources. Templates and documents read on demand, without burning context.

→ **Think of it** as the onboarding manual for a brilliant employee: already smart, just doesn't know how you work.

The best way to create a skill is by working, not by writing. The first time, don't delegate: work alongside it.



La mejor forma de crear una habilidad es trabajando, no escribiendo.

La primera vez hazlo acompañado. Al final, pídele a Houston que lo guarde.



PASO 1

Haz el trabajo con Houston

La primera vez no delegues, acompaña. Pídele el resumen, la lista, el reporte, y ve corrigiendo sobre la marcha.



PASO 2

Dale feedback en el momento

"Más directo." "Faltó el responsable." "Así no." Cada corrección que le das se vuelve una regla de la habilidad.

Lo que corriges hoy, no lo vuelves a corregir nunca.



PASO 3

Pídele que la guarde como skill

Houston escribe la habilidad solo, con todo lo que aprendió de ti en esa sesión. Tú apenas la revisas.

Queda guardada y lista para reutilizar cuando quieras.

LE DICES

"Conviértelo en una habilidad."

Tu forma de trabajar ya existe. Solo hay que grabarla una vez.

Houston · Agentes de IA que realmente hacen el trabajo

gethouston.ai

15

SLIDE 15 Do the work with it, correct it on the spot and ask it to save that. What you correct today, you never correct again.

The second best way: hand it the documents you already have. Procedure manual, onboarding guide, templates, checklists. Almost every company already has the process written down; what it doesn't have is someone who runs it the same way every time.

COPY AND PASTE, ATTACHING YOUR DOCUMENT

"I'm attaching the procedure manual we use for [the task]. Turn it into a skill of yours: pull out the step by step, the rules and the formats. Then ask me what the document doesn't say: the exceptions and the judgment calls only someone who has done this for years would know."

Piece 3 · Tools: what it uses

Without tools, an agent can only talk. It connects to the world in three ways, and each has its moment:

Un agente se conecta al mundo de tres formas. Cada una tiene su momento.



API

LA CONEXIÓN CRUDA

Le das la dirección y la llave de la app, y el agente la llama directo.

A FAVOR Todas las apps tienen una. Control total. Sirve para tus sistemas internos.

EN CONTRA Alguien la programa y la mantiene. Se rompe cuando la app cambia.

CUÁNDO Cuando no existe un conector listo para tu sistema.



MCP

EL ENCHUFE UNIVERSAL

La app se presenta sola. El agente ve qué puede hacer, sin que nadie programe.

A FAVOR La conectas una vez y sirve para cualquier agente y cualquier modelo.

EN CONTRA Alguien tiene que ofrecer el servidor. Conectar de más confunde al agente.

CUÁNDO Por defecto. Si existe, úsalo.



CLI

EL ATAJO DEL COMPUTADOR

El agente escribe comandos en la terminal, como alguien técnico de tu equipo.

A FAVOR Ya está instalado y con la sesión abierta. Rápido y muy potente.

EN CONTRA Necesita un computador con acceso. Un comando mal escrito borra cosas de verdad.

CUÁNDO Tareas técnicas o cargas masivas de datos.

EN QUÉ ORDEN ¿Existe MCP? Úsalo. ¿No hay? Mira si hay CLI. ¿Tampoco? Ahí sí, API.



En Houston le das un clic.

Houston · Agentes de IA que realmente hacen el trabajo

gethouston.ai

18

SLIDE 18 The order to follow: is there an MCP? use it. No MCP? check for a CLI. Neither? then, and only then, the API.

MCP

Default

The app introduces itself and the agent sees what it can do, **with nobody writing code**. Connect once and it works for any agent and any model.

CLI

Technical

The agent writes commands, like a technical person on your team. Fast and powerful, but **one badly written command deletes real things**.

API

Last resort

You give it the app's address and key. Good for your internal systems, but **someone has to build it and maintain it**.

In houston

One click

More than a thousand integrations already wired up under **Integrations**. Anything missing gets added from the chat as a custom integration.

The real challenge isn't connecting, it's knowing what exists

This is the actual stack the sales agent works with today:

El problema no es conectar herramientas. Es saber cuáles existen.

Este es el stack con el que trabaja mi agente de ventas hoy.

	Scraping web Apify Extrae datos de cualquier sitio, aunque no tenga API.
    	Enriquecimiento de datos Apollo · Hunter · FullEnrich · Lusha · Clay De un nombre a correo, cargo, empresa y teléfono.
	CRM HubSpot Donde viven tus contactos, tus deals y tu pipeline.
	Secuencias de outbound Instantly Manda los correos en frío y hace el seguimiento.
	Lead magnets en LinkedIn LeadShark Entrega tu contenido a quien comenta tu post.
	Secuencias en LinkedIn Waalaxy Invitaciones y mensajes en automático.
 Houston trae 1.000 herramientas pre-conectadas, con SOC 2 Type II e ISO 27001 en la capa de integración y datos.	

Houston · Agentes de IA que realmente hacen el trabajo

gethouston.ai

19

SLIDE 19 Scraping, enrichment, CRM, email sequences and LinkedIn sequences. None of this gets built from scratch: it gets connected.

 Apify · scraping	 HubSpot · CRM	 LinkedIn · source	 Gmail · email
 Slack · alerts	 Notion · tasks	 Drive · files	

Piece 4 · Memory: what it remembers

Models have a **context window**, and that is not memory: it's the desk they work on today, not the filing cabinet. It's wiped when the session closes.

Los modelos tienen algo llamado ventana de contexto, pero no es memoria.

Se borra cuando cierras la sesión. Es el escritorio donde trabaja hoy, no el archivador de lo que aprendió.



QUÉ ES

Todo lo que ve de una sola vez

Tu conversación, los documentos que le pasaste y sus instrucciones, todo junto. Se mide en tokens y hoy los modelos grandes rondan el millón.



EL ESPEJISMO

Más ventana no es mejor memoria

Un estudio probó 18 modelos y todos empeoran a medida que crece la entrada, mucho antes de tocar el límite. Pasa incluso en tareas triviales.



Y ADEMÁS

Se borra al cerrar la sesión

Nada de eso queda guardado. Sin memoria, cada conversación empieza desde cero y te toca contarle todo otra vez.

El escritorio se despeja todas las noches. Por eso el agente necesita dónde archivar.

SLIDE 21 A bigger window is not better memory: models get worse as the input grows, long before they hit the limit.

In houston, memory lives in **Settings** → **Memory**, as plain text you can read, edit, delete and take with you. To save something, tell it in the chat while you work:

COPY AND PASTE, WITH YOUR OWN RULE

"Never use em dashes in anything you write: not in chats, not in emails, not in documents. Rephrase with commas. Save this in your memory."

Piece 5 · The brain: the AI model

No model is the best at everything, and today's best isn't next week's. In houston you pick the model when you create the agent and change it whenever you want.

No necesitas 400 cuentas. Con una sola llegas a más de 400 modelos.

Hay servicios que se conectan a todos los proveedores por ti. Pagas en un solo lugar y los usas dentro de Houston.

EL QUE USAMOS



OpenRouter

UNA CUENTA, 400+ MODELOS

Recargas saldo una vez y desde ahí llamas a OpenAI, Anthropic, Gemini, DeepSeek o Mistral. Cambias de modelo sin abrir otra cuenta.



Te regalamos créditos de OpenRouter

Para que pruebes todos los modelos hoy mismo, sin poner tu tarjeta.



OpenCode

SUSCRIPCIÓN ÚNICA

Una mensualidad que te da acceso a los modelos de todos los proveedores, sin manejar saldo ni llaves por separado.

Otra alternativa, si prefieres pagar fijo.



En Houston eliges el modelo desde un menú. Le pones el mejor a cada agente y lo cambias cuando quieras.

Houston · Agentes de IA que realmente hacen el trabajo

gethouston.ai

24

SLIDE 24 You don't need 400 accounts: one OpenRouter account gets you more than 400 models. That's what you open in Step 4.



YOU SHOULD BE ABLE TO

Explain to someone, on a napkin, the five pieces of an agent and why the model is only one of them. If you can, keep going.

2

STEP 2 OF 6

Pick the right problem

~20 minutes

The most expensive mistake is starting from the technology. A good agent starts from a pain. This is the framework we used live to go from "I want an AI agent" to a specific, useful one.

First, a good candidate

The best problems for a first agent meet all four conditions:



Un buen agente empieza en un dolor, no en la tecnología.

Antes de construir, elige bien el problema. Los mejores candidatos cumplen estas 4 condiciones:



Frecuente

Pasa todas las semanas, no una vez al año.



Repetitivo

Sigue los mismos pasos cada vez que ocurre.



De bajo riesgo

Un error se corrige fácil y no te cuesta un cliente.



Medible

Puedes decir con números si funcionó.



Regla de oro: un agente nuevo no manda nada solo. Primero propone, tú revisas y apruebas.

SLIDE 27 Frequent, repetitive, low risk and measurable. Plus the golden rule: a new agent doesn't send anything on its own.

The MUUDS canvas, in five questions

Answer them in that order, on a sheet of paper or a board. After each one, **pick a single option** and continue with that.

TALLER MUUDS · RONDA 1 DE 2

Abordar un problema ambiguo con estructura permite solucionarlo con mentalidad de producto.



01

Usuario

¿Quién trabaja con el agente cada día?



02

Caso de uso

¿Qué tarea concreta resuelve?



03

Dolores

¿Qué le duele hoy a tu usuario?



04

Soluciones

¿Cómo resuelve el agente cada dolor?



05

Métricas

¿Cómo mides que funciona?

Houston

Multiplayer

Agents


Features

Pricing

Questions



- 1 **User.** Who works with the agent every day? Think roles, not departments: "a salesperson", not "sales".
- 2 **Use case.** What specific tasks does that role do? Prospect, quote, follow up, feed the CRM, book meetings.
- 3 **Pains.** What hurts today inside that use case? "Time" is not a pain: "I don't have a list of people to call" is.
- 4 **Solutions.** In how many creative ways could that pain be solved? Don't filter yet.
- 5 **Metrics.** How would you know, in numbers, that the agent is working?



That's how the canvas ended up live. That path is exactly the agent you install in Step 5.

Do it with your own agent, not on a blank page

If getting started is hard, let it interview you. After installing it, paste this into a chat:

COPY AND PASTE, WITH YOUR BUSINESS

"I run a business that does [what you do] and I want to build my first AI agent. Ask me the MUUDS canvas questions one by one: user, use case, pains, solutions and metrics. After each answer, help me pick a single option before moving to the next question. At the end, propose the specific agent I should build and what skills it would need."

✓ YOU SHOULD HAVE

One written sentence like: "My [role] loses time on [use case] because [pain]. The agent solves it by [solution] and I measure it with [metric]." Keep it: it's the brief for your own agent.



Install houston

~15 minutes

Here's where the keyboard starts. If you already had it installed, **download the latest version again**: the Agent Store you'll use in Step 5 only shows up in recent versions.

- 1 Go to gethouston.ai and click **Download app**.
- 2 When it asks for the **invite code**, type **getsh*tdone**. Don't use the waiting list button: that one leaves you waiting for an email you don't need.
- 3 Download the version for your system: **Mac**, **Windows** or **Linux**.
- 4 Install it. On Mac a window opens and you have to **drag the houston icon into the Applications folder**. On Windows you run the normal installer.
- 5 Open houston, create your account with whatever email you prefer and choose the language in **Settings**.
- 6 **Connect your AI provider**. In the left panel, **AI Models** → pick Anthropic or OpenAI → **Connect** → you sign in in the browser and come back. With a subscription, the login is all you need, no keys.
- 7 **Accept the integration permissions**. houston will ask you to authorize the provider that connects the thousand apps. Leave everything selected: if you strip permissions here, the tools fail later with no explanation.

**we're on a closed beta**

If you were invited, paste your code and download the app.

 Download for Mac

or join our waiting list to get access

By joining you agree to receive product updates from Houston.
Unsubscribe anytime. See our [Privacy Policy](#)

Want to skip the waitlist?



SCREENSHOT 1 The invite code window. Paste `getsh*tdone` and the download button activates. Don't use **Join the waiting list**.

**IF YOU'RE ON WINDOWS**

Windows may say the app is from an unknown publisher. That's normal: the certificate is recent. Click **More info** and then **Run anyway**.

Windows protegió su PC

SmartScreen de Microsoft Defender impidió el inicio de una aplicación desconocida. Si ejecutas esta aplicación, podrías poner en riesgo tu PC.

Aplicación: Houston_0.4.25_x64_en-US (1).msi


Editor: Editor desconocido

Ejecutar de todas formas

SCREENSHOT 2 The Windows warning. **More info** and then **Run anyway** appears. It only happens on Windows and only the first time.



- **The left panel** is your agents, and at the top **Mission Control**, **Integrations**, **Skills**, **AI Models**, **Agent Store** and **Settings**.
- **The center** is the mission board: **Running**, **Needs you** and what's already done. Each mission is a task you gave the agent.
- When something lands in **Needs you**, that's your agent asking for permission. That's the golden rule at work.



SCREENSHOT 3 The main screen once you're in: your agents on the left and the mission board in the center, with **Running**, **Needs you** and **Done**.

Before moving on, write to it. Any question about your business. Make the first conversation yours.

✓ YOU SHOULD SEE

houston open on the main board, with at least one provider connected under **AI Models**.
If the welcome tour got stuck, go to [Troubleshooting](#): there's a way out.



Open Apify and OpenRouter

~15 minutes

Two accounts, two pieces of the agent.



75 USD free

Eyes on the web: it pulls everyone who commented on the post, without LinkedIn blocking your account.

Email signup failing? Sign in with Google. The coupon is already in the link.

Create account ↗



10 USD free

The brain: a single account to use more than 400 models inside houston, without opening an account at every lab.

The code is redeemed under **Credits**, not on the payment screen.

Redeem credits ↗



Scan it and the coupon applies by itself.

HOUSTON_ROCKET



Scan it, create the account and paste the code under **Credits**.

OR-GETHOUSTON

Apify

- 1 Go to apify.com and create your account. If email signup gives you trouble, sign in with Google.
- 2 Apply the coupon **HOUSTON_ROCKET** to get 75 dollars in credit, valid for two months.
- 3 If you already had an account, the coupon is applied from inside: **Billing** → **Subscription** → **Add promo code**.

OpenRouter

- 1 Go to openrouter.ai and create your account. Choose **Individual** unless you're sharing it with your team.

2 If the payment screen shows up, close it. The code doesn't get applied there.

Houston

Multiplayer

Agents

Features

Pricing

Questions



3 In the left panel go to **Credits** (or open openrouter.ai/redeem) and paste

OR-GETHOUSTON, all caps.

4 Check that the balance shows up in your OpenRouter account.




PROMOTIONAL CREDITS LOOK ODD

Inside houston your OpenRouter balance may show as zero, because it only counts purchased credits. The promotional ones are there and they do work: check them on OpenRouter's **Credits** screen.

Connect OpenRouter inside houston

Seven screens and it's done. The round trip to OpenRouter starts from houston itself, so you won't get lost.


1. Find the provider. In the left panel go to **AI Models** and type **OpenRouter** in the search box.



SCREENSHOT 4 It shows up under **Providers**, with the models it brings. Click the **+** or the name.


2. Open its card and hit connect. You'll see the full catalog of models that this single account

op Houston Multiplayer Agents Features Pricing Questions 




SCREENSHOT 5 Models from every lab in one list: Anthropic, Google, OpenAI, Moonshot and the rest. Top right, **Connect**.

3. Ask for your key from here. The **Connect OpenRouter** dialog opens. Click **Get your API key**: it takes you straight to the right page on OpenRouter.




SCREENSHOT 6 Leave this window open. You'll come back to it in a minute with the key copied.

4. Create a new key. On OpenRouter you land on **API Keys** . Click **New Key** at the top right.




SCREENSHOT 7 If you already had keys, don't reuse an old one: create one for houston so you can revoke it without breaking anything else.

5. Name it and create it. A recognizable name (for example **houston**), no expiry. The credit limit must be greater than 0. The key must be created in the **API Keys** section. [link](#) [Multiplayer](#) [Agents](#) [Features](#) [Pricing](#) [Questions](#) 




SCREENSHOT 8 When you hit **Create**, OpenRouter shows you the key **once and only once**. Copy it right then.

6. Paste it into houston. Go back to the window you left open and paste the key in the **API key** field.



SCREENSHOT 9 It's masked with dots. houston stores it encrypted and the agent never sees it.

7. Done. The card switches to **Connected** and below it says **Signed in with OpenRouter**. From that moment on, any agent of yours can use those models.



SCREENSHOT 10 Connected. If one day you want to drop it, the **Sign out** button is on the same card.



THE KEY IS SHOWN ONLY ONCE
If you close the OpenRouter window without copying it, you can't get it back. Don't waste time looking: create a new key and move on. And never paste it into the chat, only into the connection field.



YOU SHOULD HAVE

An Apify account with credit and an OpenRouter account with a balance, connected under **AI Models** in houston.


5

STEP 5 OF 6

Install the **Sales Rep** agent

~10 minutes

This is the agent you saw working live. It comes with its instructions, its prospecting skill and the memories it learned on the job. You install it from the Agent Store, like an app.



STEP BY STEP Type **Felipe** in the Agent Store search box and the agent's card shows up: its skill, its tools and the **Try it now** button. In the screenshot it still has the name it had on bootcamp day; it's now called **Sales Rep**.

1

In the left panel, open **Agent Store**.

2

Type **Felipe** in the search box. The **Sales Rep** agent will show up.

3 Open it, look it over and click **Try it now** .

Houston

Multiplayer

Agents

Features

Pricing

Questions



4 houston will ask whether it should **review the agent before installing it**. Say yes: it's an


automatic security review that looks for hidden malicious instructions. Always do this with third-party agents.

5 **Continue** → give it a name. Use the role you're giving it, for example **Sales Rep**.

6 Choose **which skills to bring**. This agent shares one: LinkedIn comments to Google Sheets. Bring it.

7 Choose **which memories to bring**. Bring them all: they include which Apify actors work and mistakes already corrected. You can edit or delete them later from **Settings** → **Memory** .

8 Wait about 30 seconds while it imports.



SCREENSHOT 11 Before installing, houston summarizes what the agent brings and offers to review it. Always choose **Yes, review it**: it looks for obvious risks before letting it in.

Look inside before using it

Open the agent and go to **Settings**. You'll find the five pieces from Step 1, already filled in:

- **Instructions**: the role, the behavior and the hard rules of the commercial operator. They're in full at the end of this guide.
- **Skills**: the LinkedIn comments to Google Sheets one. Open it and read it.
- **Memory**: what it learned on the job.
- **Integrations**: empty for now. You fix that in Step 6.

✓ YOU SHOULD SEE

Your new agent in the left panel, and inside **Settings** the instructions, one skill and several memories already loaded.

6

STEP 6 OF 6

Put it to work prospecting

~30 minutes

The agent already knows what to do. Now it needs hands: Apify to read LinkedIn and Google Sheets to write the database.



One click

IN HOUSTON

Integrations → Apify MCP → Connect →

Authorize

Pick **Apify MCP**, not plain "Apify": that one asks for a token and confuses everybody.



Google Sheets

One click

IN HOUSTON

Integrations → Google Sheets →

Connect → Sign in with Google

Select the exact "Google Sheets" in the list and accept **all** the permissions.

Houston

Multiplayer

Agents

Features

IN APOLLO

Settings → Integrations → API Keys →

Copy

Only if you want the emails. Apollo shows the key **once and only once**: copy it the moment you create it.

Pricing

Questions



/API key - optional

IN HUNTER

Settings → API → Copy

Backup for the emails Apollo doesn't find. If you skip it, tell the agent to use Apollo only.




CAREFUL WITH API KEYS

Paste them under **Integrations**, **never in the chat**: there they sit exposed in the history. And don't share your screen while copying them. Lost looking for one? Ask the agent for help: *"walk me step by step through finding the API key for this tool."*

Connect the two tools

- 1 Open your agent and go to **Integrations**.
- 2 Search for **Apify**. Two results will appear. Pick the one that says **Apify MCP**, not the one that says just **Apify**.
- 3 Click **Connect** and sign in with the Apify account you opened in Step 4. It won't ask you for any key.
- 4 Search for **Google Sheets** and connect it with your Google account. Accept every permission it asks for.



SCREENSHOT 12 The two Apify results under **Integrations**. Connect the one on the right, **Apify MCP**. The one on the left is the one that asks for a token.



THE NUMBER ONE MISTAKE OF THE WHOLE BOOTCAMP

If at any point the agent asks you for an **Apify API token**, it's because you connected the wrong integration. Disconnect it and connect the one that says **MCP**. You can also just ask the agent: *"switch the Apify API for the MCP one"*.



PASO 3 · PRUÉBALO

Ahora vamos a probar el agente que instalamos.

Abre un chat con él y pégale el enlace de la publicación de LinkedIn.

1 Abre un chat con el agente

El que acabas de instalar desde la tienda.

2 Pega el enlace de la publicación

La misma foto que publicamos al principio. El link está en el chat.

3 Déjalo trabajar

Va a prospectar a quienes comentaron, igual que el mío.

Houston · Agentes de IA que realmente hacen el trabajo

ESCRÍBELE ALGO ASÍ

Prospéctame a la gente que comentó
este post:

`https://linkedin.com/posts/...`

No tiene que ser exacto. Escríbelo como se lo pedirías a una persona.

gethouston.ai


37

SLIDE 37 Open a mission, select the skill and paste the link. It doesn't have to be precise: write it the way you'd ask a person.

- 1 Get the link to a **LinkedIn post with lots of comments**. It can be yours or from anyone with a following in your industry. Copy the link to the post, not to the profile.
- 2 On the agent's board, click **New mission**.
- 3 **Select the skill** for LinkedIn comments from the list that appears. This is the most forgotten step: if you don't select it or name it, the agent improvises.
- 4 Choose the working mode. Starting out, use **Ask first**: it checks with you when it isn't sure and you give it feedback. Leave **Autopilot** for when you already trust it.
- 5 Pick the model it will use and write the instruction.

COPY AND PASTE, WITH YOUR LINK

"Prospect the people who commented on this post: `https://www.linkedin.com/posts/...`"



SCREENSHOT 13 The mission starting: the agent reads the skill, takes the post link and announces the plan before touching anything. At the bottom you can see the **Skills** selector and the model selector.

What's going to happen

- 1 **It reads its skill and creates the sheet** It sends you the Google Sheets link immediately, before finishing.
- 2 **It extracts the commenters with Apify** It goes out to LinkedIn at the network level, not browsing like a robot.
- 3 **It writes the rows in batches** Name, headline, company, comment and profile link. You watch the sheet fill up.
- 4 **It looks for the emails** If you connected an enrichment tool. This is the part that fails most.

i THIS TAKES A WHILE, AND THAT'S FINE

A real extraction over hundreds of comments takes several minutes. Let it work and go do something else. Live, with 264 people watching, the sheet filled in batches while we kept going with the workshop.



At the bootcamp the email enrichment failed: the agent reported that the responses came back too large and that it had only filled two rows. The fix wasn't redoing anything, it was answering it:

COPY AND PASTE WHEN SOMETHING COMES OUT INCOMPLETE

"The emails came out incomplete. Run the enrichment again only for the empty rows, in smaller batches, and write the results as you get them."

When it comes out right, close the loop:

COPY AND PASTE AT THE END

"Save what we learned in this task as a skill, so next time you do it this way from the start."

Optional · get the emails

To go from names to emails you need an enrichment tool. The bootcamp used **Apollo**: free account, connected by API key from [Settings](#) → [Integrations](#) → [API Keys](#). In Latin America a single tool finds between 40 and 50 percent of the emails; combining two or three (Hunter, Lusha, FullEnrich) gets you to 70 or 80.

✓ YOU SHOULD HAVE

A Google Sheet with real people, pulled from a real post, that your agent created and filled on its own. That's an agent working.

BETA PROGRAM

That was one agent and one case



In a company it's ten: one per role, with your operation's skills inside and connected to the tools your team already uses. houston is in beta and we're letting teams in in batches.

[Join the beta program ↗](#)



You've got the method and a working agent. Now point it at something that actually hurts.

Start with the boring, not the important

It sounds backwards and it isn't. The first task you give a new intern isn't the negotiation with your biggest customer. Pick something repetitive, that steals your time, and where a mistake doesn't cost much.

Create your own agent in three sentences

You don't need to know how to assemble anything. Create a new agent, name it and paste this into the chat:

COPY AND PASTE, WITH YOUR ROLE

"You're going to be my [bookkeeping assistant / purchasing analyst / whatever you need]. Research how that role works and write your own instructions: role, behavior, hard rules, vocabulary and tone. Then propose 10 skills we could build together, ordered by how much time they'd save me."

Change the source, not the method

The agent you installed prospects from LinkedIn, but the pattern works the same on other sources. Ask it like this:

- **Instagram or TikTok.** "I want to do the same thing but with the comments on this Instagram account. What do you need?"
- **Google Maps.** "Use Apify to pull the restaurants in this area with their phone number and website."
- **Events.** "Take this attendee list, find them on LinkedIn and build the database."
- **Public data.** Chambers of commerce, public tenders and government registries.

Five ideas outside of sales

- **The weekly report.** The one you build by hand on Fridays and hate.
- **The same answers as always.** The emails you reply to identically every week.
- **Quote follow-ups.** Who hasn't replied and for how long.
- **Document handling.** Collecting, merging and filing the PDFs for each process.
- **Meeting prep.** Have it build the customer context before you walk in.

Four tweaks worth gold

Everything is edited by talking to it in plain language, no code. These four came out of real use:

1. **Have it qualify before spending credits.**

COPY AND PASTE, WITH YOUR CRITERIA

Houston

Multiplayer

Agents

Features

Pricing

Questions



"Modify the skill: create a column called qualification and always qualify against these ideal customer criteria: [your criteria]. If a prospect isn't qualified, don't even fetch the email."

2. Have it use the skill without you reminding it.

COPY AND PASTE INTO THE CHAT

"Save in your memory that every time I send you a LinkedIn post link, you use the LinkedIn comments to Google Sheets skill."

3. Have it stop repeating prospects across runs.

COPY AND PASTE INTO THE CHAT

"Whenever we're prospecting, check that we're not enriching duplicate data from previous runs."

4. Have it connect to a tool it doesn't know. This works with any app that has public documentation.

COPY AND PASTE, WITH YOUR TOOL

"I want to connect to [your tool], please help me. Here's the link to its documentation: [paste the link]"

The next level: the full campaign

The prospect sheet is half the road. With two more tools, the same agent builds the campaign and leaves it **paused**, waiting for your approval:



API key

Finds verified work emails. Free up to ~100 a month. In Latin America it finds between 40 and 60 percent.



API key

Sends the email sequence without burning your domain. For volume, use warmed-up alternate domains.



BEFORE YOU ACTUALLY SEND

Loading the campaign is safe: it stays paused. But don't activate it until you have properly configured sending inboxes. If you don't yet, tell the agent: *"I want to test the skill without sending anything for now."*



an event: *"every Monday at 8, check the CRM and send me the deals with no movement"*. First do the task well once by hand; only then schedule it.



ONE SINGLE NEXT STEP

This week: **one boring task, one agent**. Just one. Whoever does one thing comes back; whoever leaves with ten ideas does none.

The full instructions of the Sales Rep agent

These are, word for word, the instructions of the agent you installed. Read them as a template: notice the order (role, how it talks to you, skills, protocol, transport, data rules and what it never does) and replace the content with your own business.

1

Role. What it covers and what it doesn't. "I draft. I never send."

2

How I talk to you. Bans technical jargon in the chat and gives the translation table.

3

18 skills. Grouped by moment in the process, not by tool.

4

What I never do. The hard rules. It's the most important section.

FULL INSTRUCTIONS · COPY THEM AND ADAPT THEM

Role

I'm your full commercial operator. Prospecting, pipeline, meetings, proposals, CRM, retention and expansion · all in one conversation, one context record, one deliverables folder.

I draft. I never send. I never move CRM stages or change prices without your approval. You decide.

At the end of every conversation, if you tell me DONE, I read the whole session and save whatever is worth keeping as a skill or a memory.

To get started

No upfront onboarding. Tell me what you want to do and I get to work. When I need something specific (company, ideal customer, pricing, pipeline stages, qualification framework, connected CRM) I ask one precise question, save the answer in your sales context and move on.

Ways to give me context, in order of preference: connected app > file > URL > pasted text. Connect



How I talk to you

You're not technical. You don't care about file names, paths or JSON. In the chat I never say internal file names, paths, plumbing words (schema, JSON, config, manifest) or internal tools.

I talk about things as what they are to you:

| I don't say | I say |

| "Updating your `idealCustomer` in `sales-context.md`" | "Updating your ideal customer profile" |

| "writing to `context-ledger.json`" | "saving this in your sales context" |

| "added a skill at `.agents/skills/foo/SKILL.md`" | "created a new skill called Foo" |

| "logged in `leads.json`" | "added it to your prospect list" |

| "notes in `calls/{slug}/`" | "saved the call notes" |

| "the `outputs.json` index" | "your saved work" |

| "saved in `learnings.json`" | "I'll remember it" |

Internally I still read and write those files · that doesn't change. The rule is only for what shows up in the chat.

Exception: if you use a technical term first, I answer in that same register.

My skills (18 in total)

Setup

- `set-up-my-sales-info` · tell me about your company, ideal customer, pricing, pipeline stages and how you handle objections. The base every other skill uses.

- `profile-my-buyer` · precise profile of who buys in a segment: champion, economic buyer, blocker, disqualifiers, anchor accounts.

Daily rhythm

- `brief-me-for-today` · one-screen morning summary: today's meetings, drafts pending approval, 3 key moves, follow-up list.

Outbound prospecting

- find-me-leads · fresh prospects in a segment from the source you choose (CRM lookalikes, LinkedIn thread, funding feed, Google Maps, subreddit). Scored against the playbook, trigger signal cited.
- research-an-account · depending on the depth: quick-qualify (fast URL read) | full-brief (cited multi-source research) | enrich-contact (specific person enrichment) | warm-paths (first-degree intros via LinkedIn + CRM).
- score-my-pipeline · depending on the subject: lead (score unscored prospects) | lead-fit (fit + angle for one prospect) | deal-health (all open deals) | customer-health (all current customers). Top 2 drivers per row.
- write-my-outreach · depending on the stage: cold-email | cold-script | followup | inbound-reply | renewal | churn-save. Tone matched to your voice, drafts only.

Meetings and deal work

- prep-a-meeting · depending on the type: call (pre-call one-pager with questions for the weakest qualification pillar) | account-review (quarterly review pack with delivered outcomes, usage trend, risks, renewal runway).
- capture-my-call-notes · transcript or recording → structured notes (actual vs. planned agenda, attendees, verbatim pains, decisions, next steps).
- build-a-battlecard · prospect vs. competitor card: 3-criteria grid, 3 trap questions, 3 rebuttals, 2 proof points. Every claim cites a source.
- handle-an-objection · 3-sentence reframe on the call (acknowledge → anchor account example → next step with a date) plus a short post-call email in your voice.
- write-a-proposal · one-page proposal built on the deal's verbatim problem and success metric. Pricing stays inside the playbook; exceptions get flagged.
- write-a-close-plan · mutual action plan across procurement, security, budget and legal with owners (yours + theirs) and dates. Top 3 risks and unknowns up front.

CRM and forecasting

- manage-my-crm · depending on the action: clean (hygiene diff, no mutations) | query (natural



- `run-my-forecast` · classify open deals against the playbook's exit criteria into Commit / Best / Pipeline / Omit, sum annual revenue per bucket, flag week-over-week slippage.

Retention and expansion

- `plan-an-onboarding` · kickoff agenda + agreed success metric + 90-day timeline to first value. The anchor for every later account review and renewal.

- `find-my-expansions` · scan GREEN customers for usage spikes / team growth / feature requests, rank by annual revenue potential over effort.

Context protocol

Before any substantive work I read your sales context. For every required field that's missing, I ask one question with the best modality (connected app > file > URL > paste), write the answer atomically and continue. I never ask the same thing twice.

Fields the context tracks:

- `universal.company` · name, website, 30-second pitch, stage.

- `universal.voice` · voice samples + tone notes (the base for every outreach draft).

- `universal.playbook` · whether the sales playbook exists; path; last updated date.

- `universal.idealCustomer` · industry, roles, pains, triggers, disqualifiers.

- `domains.outbound` · connected prospecting sources, cadence windows, primary channel.

- `domains.crm` · connected CRM, stage names, owner map, routing policy.

- `domains.meetings` · connected call recorder, primary first-call objective, qualification framework (MEDDPICC/BANT/custom).

- `domains.retention` · connected billing and product usage sources, health thresholds, renewal notice window.

Combined flows

Some requests cross domains. Everything in one agent = I chain skills without handoffs:

- **New deal, zero to close:** `find-me-leads` → `research-an-account depth=full-brief` → `score-my-

pipeline subject=lead-fit` → `write-my-outreach stage=cold-email` → `prep-a-meeting type=call` →

multiplayer` → `check-my-pipeline subject=lead-fit` → `write-my-outreach stage=followup` → `write-a-proposal` → `write-a-close-plan`

Houston

Agents

Features

Pricing

Questions



- **Monday review:** `check-my-sales subject=sales-health` → group by domain, flag stalled deals + missed follow-ups, recommend 3 key moves.

- **Pre-call pipeline:** `build-a-battlecard` and `prep-a-meeting type=call` are built on the same `research-an-account` artifact.

- **Customer health loop:** `score-my-pipeline subject=customer-health` → red → `write-my-outreach stage=churn-save`; green with a usage spike → `find-my-expansions`.

Transport

Composio

Every connected-app tool goes through Composio. I discover the slugs at runtime with `composio search <category>` and execute by slug. If a connection is missing, I tell you which category to connect and I stop. No hardcoded tool names. Categories I use:

- **CRM** · HubSpot, Salesforce, Attio, Pipedrive, Close

- **Email** · Gmail, Outlook

- **Meetings** · Gong, Fireflies

- **Calendar** · Google Calendar, Outlook

- **Search / research** · Exa, Perplexity

- **Scraping** · Firecrawl

- **Social** · LinkedIn

- **Task management** · Linear, Notion, Asana

- **Billing** · Stripe

- **Messaging** · Slack

Kernel (cloud browser)

I use Kernel when the task needs a real browser: sites with no API, JS-heavy pages Firecrawl can't render, sites with anti-bot protection, or tasks that require clicking, filling forms or acting on a website.

- The site blocks scrapers or requires JS execution
- The task involves clicking, submitting forms or signing in
- LinkedIn or other closed sites where no Composio connector is available
- Firecrawl returns empty or blocked content

How to use it: load the `kernel-agent-browser` skill, then use `agent-browser -p kernel` commands. Always configure before the first command. The `KERNEL\_API\_KEY` value is NOT stored here · it lives in saved memory. Resolve it from memory at runtime and export it, then set the rest:

```
```bash

export KERNEL_API_KEY="<resolve from saved memory>"

export KERNEL_STEALTH=true

export KERNEL_TIMEOUT_SECONDS=300

```
```

Use `KERNEL\_PROFILE\_NAME=<site>` to persist sessions across conversations (sign in once, reuse forever). Always close the session with `agent-browser -p kernel close` when done, to save the profile.

Decision rule:

- Has a Composio connector and needs no browser action → use Composio
- No API, heavy JS, anti-bot protection, or requires a click/form → use Kernel
- Firecrawl returned empty/blocked content → switch to Kernel automatically

Data rules

- My data lives at the agent root · never under `.houston/<agent-path>/`.
- `config/` · what I've learned about you. Filled at runtime by progressive capture.
- `context/sales-context.md` · the sales playbook.
- Artifact folders at the root: `personas/`, `battlecards/`, `leads/`, `leads.json`, `accounts/`, `calls/`,

- `outputs.json` at the root indexes every artifact with `{id, type, title, summary, path, status, createdAt, updatedAt, domain}`. Atomic writes: temp file + rename. Read-merge-write, never overwrite.

- Every record carries `id` (uuid v4), `createdAt`, `updatedAt`.

What I never do

- Send, publish or commit CRM stage changes without your approval.

- Make pricing promises outside the playbook range.

- Invent customer quotes, call facts, competitor moves or pipeline numbers · if the source is weak, I mark it TBD and ask.

- Guess your ideal customer, qualification framework or pricing · I read the playbook or ask you to run `set-up-my-sales-info` first.

- Use guilt, false scarcity or dark patterns in churn-save, renewal or expansion drafts.

- Write to `.houston/<agent-path>/` at runtime.

- Hardcode tool names in skills · Composio discovery at runtime, always.

DON'T COPY IT ALL AT ONCE

This agent has months of use behind it. For yours, start with **Role**, **How I talk to you** and **What I never do**. The skills show up on their own as you work with it.

TALK TO THE TEAM

Not sure where to start in your company?



In 30 minutes we look at the process costing you the most time and tell you whether it's a good case for an agent and how we'd build it. If it isn't, we tell you that too.

Book a call ↗



You got here with one agent and one case. This section is for when you want your whole operation to work like that.

WHAT YOU JUST DID

One agent, one case, your laptop

A process that used to take you hours now runs on its own. You built it yourself, in an afternoon, without coding.



WHAT WE DO WITH TEAMS

One agent per role, with your way of working inside

Sales, support, content, back office. With your company's skills, connected to your tools and running on their own when they should.

What working with us includes

We don't sell licenses and leave you alone. The value isn't in the software, it's in the agents ending up with your operation inside them.

Week 1

Process diagnosis

We sit down with your team, look at where the time goes and pick the two or three processes actually worth it. Starting with the wrong process is the most expensive way to try this.

What you get

Agents built on your operation

No generic templates: instructions, skills and memory carrying the way your company works, connected to the tools your team already uses.

With your people

Team training

The people who'll use them learn to give them work, correct them and turn that correction into a skill. Without this, any agent dies in two weeks.

After

Support while it beds in

The first months are about adjusting. We stay until the agent is producing results your team actually uses, not a pretty demo.



It's for you if

- ✓ You already know which process is costing you hours every week.
- ✓ You have volume: sales, support, content, quotes, back office.
- ✓ You want this running in weeks, not a six-month pilot.
- ✓ Someone on the team is willing to own the agent. They don't have to be technical.

It's not for you if

- ✗ You're looking for a chatbot to answer questions on your website.
- ✗ The process doesn't exist yet: first you define it, then you automate it.
- ✗ You want the agent writing to customers with nobody reviewing.
- ✗ You expect savings without changing anything about how the team works.

How it starts

- 1 A **30 minute** call. You tell us the process costing you the most time.
- 2 We tell you straight whether it's a good case for an agent, how we'd build it and what we'd need from your side.
- 3 If it makes sense for both of us, we build the kickoff plan with your team. If not, you still keep everything in this guide.

Want this running in your company?

Book the call and we'll look at your specific case. And if for now you'd rather explore on your own, join the beta program: houston is still in beta, you get into the early access batches, and what you tell us goes straight into what we build.

[Book a call ↗](#)

[Join the beta program ↗](#)

calendly.com/gethouston/30min · forms.gle/dysscJx284wMqBX46



Houston

AND IF YOU'D RATHER KEEP GOING ALONE, PERFECT

Everything in this guide works without talking to anyone. The agent you installed is yours and keeps working the same. This section is for when the problem stops being yours and becomes the whole team's.

[Multiplayer](#)[Agents](#)[Features](#)[Pricing](#)[Questions](#)

WHEN SOMETHING BREAKS

Troubleshooting

All of this actually happened, live, with more than 500 people at once. If it happens to you, it isn't you.



Windows says the app is from an unknown publisher

It's normal and it isn't a virus: the signing certificate is recent. Click [More info](#) and then [Run anyway](#). It only happens the first time.



I downloaded houston on Mac but it isn't in Applications

When you open the downloaded file a window appears with the houston icon and a folder. You have to **drag the icon into the Applications folder**. If it asks, choose replace.



It asks for an invite code and I don't have one

Use [getsh*tdone](#) in the code field. If instead you clicked the waiting list option, you'll be waiting for an email you don't need: go back to the download page and use the code.



I got stuck in the welcome tour

This happens when you didn't grant all the permissions and the tour asks for an action the agent can't do. Go to [Integrations](#), find the half-connected integration, click [Disconnect](#) and connect it again accepting all permissions. Recent versions also offer an option to skip the tour.



The most common mistake of all. You connected plain **Apify** instead of **Apify MCP**. Disconnect it and connect the one that says **MCP**, which only asks you to sign in. You can also just write to it: *"switch the Apify API for the MCP one"*.

? The Apify coupon won't apply because I already had an account

Apply it from inside: **Billing** → **Subscription** → **Add promo code** and paste **HOUSTON_ROCKET**.

? The OpenRouter code doesn't work or says I hit the limit

Two things. First: the code doesn't go on the payment screen, it goes under **Credits** or at openrouter.ai/redeem. Second: promo codes have limited spots and can run out. If it no longer applies, OpenRouter still works by topping up your balance, or you can use your ChatGPT or Claude account directly.

? I redeemed the OpenRouter credits but houston shows zero

houston counts purchased credits, not promotional ones. The balance is there: check it on OpenRouter's **Credits** screen. The models work the same.

? I can't see the Agent Store

You have an old version of the app. Go back to gethouston.ai, download the latest version and install it over the old one.

? The agent doesn't answer or says it has no model

It has no engine. Go to **AI Models** in the left panel and connect at least one provider: Anthropic or OpenAI with your subscription (login only), or OpenRouter with an API key.



It tells you you don't have admin permissions. In Apollo go to **Settings** → **Users & Teams**, edit your profile and make yourself an admin of the workspace. Then go back to **Integrations** → **API Keys**.

? The LinkedIn extraction returns nobody

It's almost always the post: too few comments, private, or the link came out incomplete. Try another post with plenty of comments and copy the full link to the post, not to the profile.

? I asked for something and it took forever or froze

Tasks that go out to the internet take several minutes, that's normal. If it really is stuck, ask it to try again or to split the task into smaller batches.

? A Google integration rejects the action over permissions

Go to **Integrations**, find the failing one, click **Disconnect** and connect it again accepting every permission it asks for.

? WhatsApp Business won't connect

It's the trickiest integration, and that's on Meta, not Houston. You need a Meta Business account, the verified number and your WhatsApp Business account identifier (WABA ID). Ask your agent to walk you through it step by step for your case.

? I have accounts with different emails in each tool

Not a problem: your Houston account is only for signing in, and each integration connects with whatever account you want, even several at once. What is worth doing is separating workspaces, one personal and one for the company, so the agent doesn't hesitate about which account to use.



The personal space is for one person. Create a **team workspace**, and from **Settings** → **Administration** → **People** invite your colleagues. Everyone lands on the same board and can pick up each other's missions.

? I don't know how to do something inside houston

Ask the agent, in plain language: *"I have no idea what an API key is, where do I find it?"* or *"walk me step by step through connecting this tool"*. It's built to guide you, not the other way around.

? I copied an API key wrong and now I can't see it

Almost every tool shows the key **once and only once**, for security. Don't go looking: create a new key and paste it into houston.

? I switched accounts in a tool but houston still uses the old one

The previous connection is stored at the integrations provider. Go to composio.dev with the same email you used to connect your apps, find the app in your connections, delete the old one and connect it again from houston.

? Halfway through a mission it says Google Sheets isn't connected

Go to **Integrations**, look for **Google Sheets** (the exact one), connect it with your Google account and go back to the mission. Then write: *"Google Sheets is connected now, continue."* No need to start over.

? It says "LinkedIn is blocking Apify" or something similar

Don't take it at face value: sometimes the agent gets the diagnosis wrong. Write to it: *"please verify what you're telling me, or tell me if you're having a different problem."* In the live session the real cause was something else, and it found it itself once it checked.



That's expected: normal coverage is between **50 and 70 percent**, and in Latin America it's usually less. In one session, 42 commenters produced 24 emails. These are databases and people change jobs: an empty cell isn't a failure. Also, every search burns credits, so don't re-run the missing ones without a reason.

? I'm about to update houston and I'm afraid of losing everything

Relax. Your agents, skills and memories live in your account, not inside the installer. You can uninstall and reinstall: when you sign in, it's all still there.

? I lost track of which step I was on

Use the success checklist. Tick what you've already done and continue from the first unticked item. And if you're truly stuck, write in the bootcamp WhatsApp group: we're there.



None of this solved it?

Post it in the bootcamp WhatsApp group. The team is there and so are the people who already went through the same thing. Nobody is left without their agent running.

[Join the community ↗](#)

chat.whatsapp.com/BdLpn8T92DL4pRahiqOZE

BEFORE YOU CLOSE

Your **success** checklist

If you can tick all of this, you're done.



- ☐ I wrote my canvas sentence: user, use case, pain, solution and metric
- ☐ houston installed, with an **AI model** connected
- ☐ **Apify** account created
- ☐ **OpenRouter** account created and connected
- ☐ **Sales Rep** agent installed from the Agent Store
- ☐ I read its **instructions** and its **skill** from the inside
- ☐ **Apify MCP** and **Google Sheets** connected
- ☐ I have a Google Sheet with **real prospects** from a real post
- ☐ I gave it **feedback** at least once and asked it to save that
- ☐ I picked **one boring task** in my business for this week



Esto es todo lo que hiciste hoy.

TEORÍA

Lo que ahora entiendes

- 01 Qué es un agente**
Chatbot, co-worker y agente no son lo mismo.
- 02 Las 5 piezas**
Identidad, habilidades, herramientas, memoria y el modelo.
- 03 Instrucciones que funcionan**
Rol, comportamiento, reglas duras, vocabulario y tono.
- 04 Cómo se crean las habilidades**
Trabajando con el agente, o subiéndole tus documentos.
- 05 Herramientas y memoria**
Con qué se conecta y cómo deja de olvidar.

PRÁCTICA

Lo que construiste

- 06 Diseñaste tu agente**
Con el canvas MUUDS, partiendo de un dolor real.
- 07 Instalaste las herramientas**
Houston, Apify y OpenRouter, con tus créditos.
- 08 Instalaste tu primer agente**
Desde la tienda de agentes de Houston.
- 09 Lo pusiste a trabajar**
Prospectó en vivo desde una publicación real.



Entraste sin saber qué era un agente. Sales con uno funcionando y con el método para construir el siguiente.

Houston · Agentes de IA que realmente hacen el trabajo

gethouston.ai

38

SLIDE 38 You came in not knowing what an agent was. You leave with one working and with the method to build the next one.



Tell us what you built

The bootcamp WhatsApp group is still alive: that's where the codes, the new tutorials and everyone's questions go. Scan the code or use the button, and tell us what agent you put to work.

[Join the community ↗](#)

chat.whatsapp.com/BdLpn8T92DL4pRahiqOZEer

And if you want it for your whole team

You already know how to build one. When the process stops being yours and becomes the whole team's, we set it up with you: diagnosis, agents with your operation inside them, and training for the people who'll use them. It's all in [houston for your team](#).

[Book a call ↗](#)

[Join the beta program ↗](#)

Stuck?

Houston

[Multiplayer](#)

[Agents](#)

[Features](#)

[Pricing](#)

[Questions](#)



Don't stay there. Write in the WhatsApp group or email me at felipe@gethouston.ai and we'll

get you moving. Nothing that happens to you is strange: everything in this guide happened to us first, live, with 500 people watching.

And remember the one next step: this week, one boring task, one agent.

Houston

The workspace where people and AI agents work together. Free to try, for the whole team when you are ready.

[Download](#)

[macOS](#) · [Windows](#)

Company

[Vision](#)

[Contact](#)

Product

[Multiplayer](#)

[Agents](#)

[Features](#)

[Pricing](#)

[Questions](#)

Resources

[Agent Store](#)

[Guides](#)

[Changelog](#)

[GitHub](#)

[Twitter / X](#)

Houston. Open source.

[Privacy Policy](#)

[Terms of Service](#)

[Unsubscribe](#)